

Kernel Module Project

An overview of some of the implementation details

Pranav Batra

IPC via /proc VFS files

- Read and write messages are marshalled into the appropriate structures. These structures allow for block-based interprocess communication between the userland program and the kernel module instead of a byte-oriented data stream.
- These same structures can be used with bidirectional multicast socket-based communication protocols (such as netlink or netlink-generic) instead of /proc files.
- Proc files are easier to work with

```
enum action {  
    NOP, BLOCK, LOG, ENABLE_LOG_BLOCK, DISABLE_LOG_BLOCK, RESET_ALL  
};  
  
enum param {  
    BLACK, WHITE, ADD_ADDR, REMOVE_ADDR, REMOVE_ALL, RESET  
};  
  
struct writemsg { //pod struct  
    enum action action;  
    enum param param;  
    char addr[16];  
};  
  
struct readmsg { //pod struct  
    char saddr[16];  
    char daddr[16];  
    int hook;  
    bool blocked;  
    struct timespec t;  
};
```

Build tools

- Kbuild is the build system for used to compile the Linux kernel.
- Supports in-place builds however out-of-source builds are trickier
 - One solution involves using hardlinks
 - Clion only supports out-of-source compilation
- CMake is used to generate the makefile
- Objcopy converts the compiled module into an object file
 - Necessary to embed the module into the userland executable.

```
cmake_minimum_required(VERSION 2.6)
set(CMAKE_RUNTIME_OUTPUT_DIRECTORY "${CMAKE_SOURCE_DIR}/bin")
set(CMAKE_C_FLAGS "-Wall")
SET(CMAKE_CXX_FLAGS "${CMAKE_CXX_FLAGS} --std=c++11")
add_executable(main main.cpp ${CMAKE_RUNTIME_OUTPUT_DIRECTORY}/netmod.ko.o)
add_custom_command(
    OUTPUT ${CMAKE_RUNTIME_OUTPUT_DIRECTORY}/netmod.ko.o
    COMMAND ./build_module.sh
    WORKING_DIRECTORY ${CMAKE_SOURCE_DIR}
    DEPENDS netmod.c header.h
)
#!/bin/sh
mkdir -p bin
ln -f netmod.c bin/netmod.c
cat > bin/Makefile << "EOF"
obj-m += netmod.o
ccflags-y := -I$(src)/../
all:
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) modules
clean:
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) clean
EOF
cd bin && make 2>&1 | sed -r 's,^([^\:]*\/)bin/(.\/)?,\1,'
rm -f Makefile netmod.c
objcopy -I binary -O elf64-x86-64 -B i386 netmod.ko netmod.ko.o
```

Init_module, cleanup_module

- Called when the module is loaded and unloaded, respectively
- Register two netfilter hooks, NF_INET_LOCAL_IN for inbound packets and NF_INET_LOCAL_OUT for outbound ones.
- Set up the proc file with permissions 0666 – everyone is allowed to read from & write to the vfs file
- The name of the proc file is given in header.h: netmod_hello

```
int init_module(void) {  
    in = template;  
    in.hooknum = NF_INET_LOCAL_IN;  
    in.priority = NF_IP_PRI_FIRST;  
    nf_register_hook(&in);  
    out = template;  
    out.hooknum = NF_INET_LOCAL_OUT;  
    out.priority = NF_IP_PRI_LAST;  
    nf_register_hook(&out);  
    proc_create(PROC_NAME, 0666, NULL, &proc_fops);  
    printk("module loaded (%s)\n", path);  
    return 0;  
}
```

```
void cleanup_module(void) {  
    nf_unregister_hook(&in);  
    nf_unregister_hook(&out);  
    remove_proc_entry(PROC_NAME, NULL);  
    printk("module unloaded (%s)\n", path);  
}
```

Ring buffers

- A ring buffer of readmsg's is used to store information about the logged packets
 - When the buffer is full, messages are overwritten in FIFO order – i.e., the oldest message are dropped first.
 - The index keep looping around, hence the name.
- Better to drop packets than to slow the system to a crawl!
 - This is what happens when you write to the log file directly from the kernel, my initial approach to logging.
- Bufsize is defined as 256 in header.h

```
if (log) {
    rbuf[ri].hook = state->hook;
    rbuf[ri].blocked = !allow;
    getnstimeofday(&rbuf[ri].t);
    snprintf(rbuf[ri].saddr, 16, "%pI4", &iph->saddr);
    snprintf(rbuf[ri].daddr, 16, "%pI4", &iph->daddr);
    ri++;
    ri = ri % bufsize;
    rl = min(rl + 1, bufsize);
}

rl--;
ri--;
if (ri < 0)
    ri += bufsize;
copy_to_user(buf, &rbuf[ri], count);
```

That's all folks

- There is also a userland program that interfaces with the kernel module. It does some interesting stuff with regards to file descriptor manipulation and the use of fork/exec in order to overlay the kernel log file and the list of logged packets atop the CLI (command line interface).
- Take a look at the github source code for more information about the internals of both the userland program and the kernel module. Comments are included in the code!
- <https://github.com/CS3281-2016/project>